

# Week 1: R basics

Katie Schuler

## What is data science?

Data are descriptions of the world around us, collected through observation and stored on computers. Computers enable us to infer properties of the world from these descriptions. Data science is the discipline of drawing conclusions from data using computation.

– *Computational and Inferential Thinking: The Foundations of Data Science*

Data science is about making decisions based on incomplete information. This concept is not new – brains (especially human brains!) love doing this. While the underlying concept is not new, the computational tools we use are *relatively* new. And we have a lot more data!




Figure 1: from <https://web-assets.domo.com/miyagi/images/product/product-feature-22-data-never-sleeps-10.png>

## Data science workflow

The folks who wrote [R for Data Science](#) proposed the following data science workflow:

Let's unpack what these mean briefly:




Figure 2: from R for Data Science

1. **Import:** gather data from a variety of sources, which can include structured data (like databases and spreadsheets) and unstructured data (like text, images, and videos).
2. **Tidy and Transform:** The raw data we import is often messy. Data scientists clean and preprocess the data, which involves removing errors, handling missing values, and transforming data into a suitable format for analysis.
3. **Visualize (exploratory data analysis)** visualize and summarize data to identify patterns, form hypotheses, select appropriate models, and guide further analysis.
4. **Model:** Using statistical methods, machine learning algorithms, and other computational techniques, data scientists build models to understand underlying patterns in the data. Models are tested using validation techniques to ensure their accuracy and reliability. Then data scientists use them to draw meaningful conclusions, like predictions about the future or inferences about population.
5. **Communicate:** Finally, a crucial part of data science is communicating findings clearly and effectively, whatever your purpose (academic, industry, or the public!)
6. **Program:** Surrounding all these is programming, since the computational tools are what make these possible!

Data science has applications in many fields far beyond language and the mind. It allows us to make data-driven decisions, solve complex problems, and uncover hidden insights that might not be apparent through other methods.

## Overview of the course

We will spend the first few weeks getting comfortable **programming in R**, including some useful skills for data science:

- R basics
- Data visualization
- Data wrangling

Then, we will spend the next several weeks building a **foundation in basic statistics and model building**:

- Sampling distribution
- Hypothesis testing
- Model specification
- Model fitting
- Model evaluation

Finally we will cover a selection of **more advanced topics** that are often applied in language and mind fields, with a focus on basic understanding:

- Feature engineering
- Classification
- Mixed-effect models

## Why R?

With many programming languages available for data science (e.g. R, Python, Julia, MATLAB), why use R?

- Built for stats, specifically
- Makes nice visualizations
- Lots of people are doing it, especially in academia
- Easier for beginners to understand
- Free and open source (though so are Python and Julia, MATLAB costs \$)

If you are interested, here is a [math professor's take](#) on the differences between Python, Julia, and MATLAB. Note that although they're optimized for different things, **they are all great** and the technical skills and conceptual knowledge you gain in this course will transfer to other languages.

## Google Colab

There are many ways to program with R. Some popular options include:

- [R Studio](#)
- [Jupyter](#)
- [VS Code](#)
- and even simply the [command line/terminal](#)

**Google Colab** is a cloud-based Jupyter notebook that allows you to write, execute, and share code like a google doc. We use Google Colab because it's simple and accessible to everyone. You can start programming right away, no setup required! Google Colab officially supports Python and R (and secretly [Julia](#), too!)

New R notebook:

- [colab \(r kernel\)](#) - use this link to start a new R notebook
- `File > New notebook` and then `Runtime > Change runtime type` to R

Cell types:

- `+ Code` - write and execute code
- `+ Text` - write text blocks in [markdown](#)

Left sidebar:

- `Table of contents` - outline from text headings
- `Find and replace` - find and/or replace
- `Files` - upload files to cloud session

Frequently used menu options:

- `File > Locate in Drive` - where in your Google Drive?
- `File > Save` - saves
- `File > Revision history` - history of changes you made
- `File > Download > Download .ipynb` - used to submit assignments!
- `File > Print` - prints
- `Runtime > Run all` - run all cells
- `Runtime > Run before` - run all cells before current active cell
- `Runtime > Restart and run all` - restart runtime, then run all

Frequently used keyboard shortcuts:

- `Cmd/Ctrl+S` - save
- `Cmd/Ctrl+Enter` - run focused cell
- `Cmd/Ctrl+Shift+A` - select all cells

- Cmd/Ctrl+/ - comment/uncomment selection
- Cmd/Ctrl+] - increase indent
- Cmd/Ctrl+[ - decrease indent

## R Basics

### Basics

We begin by defining some basic concepts:

- **Expressions** are combinations of values, variables, operators, and functions that can be evaluated to produce a result. Expressions can be as simple as a single value or more complex involving calculations, comparisons, and function calls. They are the fundamental building blocks of programming.
  - 10 - a simple value expression that evaluates to 10.
  - `x <- 10` - an expression that assigns the value of 10 to `x`.
  - `x + 10` - an expression that adds the value of `x` to 10.
  - `a <- x + 10` - an expression that adds the value of `x` to 10 and assigns the result to the variable `a`
- **Objects** allow us to store various types of data, such as numbers, text, vectors, matrices; and more complex structures like functions and data frames. Objects are created by assigning values to variable names with the assignment operator, `<-`. For example, in `x <- 10`, `x` is an object assigned to the value 10.
- **Names** that we assign to objects must include only letters, numbers, `.`, or `_`. Names must start with a letter (or `.` if not followed by a number).
- **Attributes** allow you to attach arbitrary metadata to an object. For example, adding a `dim` (dimension) attribute to a vector allows it to behave like a matrix or n dimensional array.
- **Functions** (or commands) are reusable pieces of code that take some input, perform some task or computation, and return an output. Many functions are built-in to base R (see below!), others can be part of packages or even defined by you. Functions are objects!
- **Environment** is the collection of all the objects (functions, variables etc.) we defined in the current R session.
- **Packages** are collections of functions, data, and documentation bundled together in R. They enhance R's capabilities by introducing new functions and specialized data structures. Packages need to be installed and loaded before you can use their functions or data.
- **Comments** are notes you leave to yourself (within code blocks in colab) to document your code; comments are not evaluated.

- **Messages** are notes R leaves for you, after you run your code. Messages can be simply for-your-information, warnings that something unexpected might happen, or errors if R cannot evaluate your code.

Ways to get help when coding in R:

- **Read package docs** - packages usually come with extensive documentation and examples. Reading the docs is one of the best ways to figure things out. Here is an example from the [dplyr package](#).
- **Read error messages** - read any error messages you receive while coding — they give clues about what is going wrong!
- **Ask R** - Use R's built-in functions to get help as you code
- **Ask on Ed** - ask questions on our class discussion board!
- **Ask Google or Stack Overflow** - It is a normal and important skill (not cheating) to google things while coding and learning to code! Use keywords and package names to ensure your solutions are course-relevant.
- **Ask ChatGPT** - You can similarly use ChatGPT or other LLMs as a resource. But keep in mind they may provide a solution that is wrong or not relevant to what we are learning in this course.

## Important functions

For objects:

- `str(x)` - returns summary of object's structure
- `typeof(x)` - returns object's data type
- `length(x)` - returns object's length
- `attributes(x)` - returns list of object's attributes
- `x` - returns object x
- `print(x)` - prints object x

For environment:

- `ls()` - list all variables in environment
- `rm(x)` - remove x variable from environment
- `rm(list = ls())` - remove all variables from environment

For packages:

- `install.packages()` to install packages
- `library()` to load the package into your current R session.
- `data()` to load data from package into environment
- `sessionInfo()` - version information for current R session and packages

For help:

- `?mean` - get help with a function
- `help('mean')` - search help files for word or phrase
- `help(package='tidyverse')` - find help for a package

## Vectors

One of the most fundamental data structures in R is the **vector**. There are two types:

- **atomic vector** - elements of the same data type
- **list** - elements refer to any object (even complex objects or other lists)

Atomic vectors can be one of six **data types**:

- **double** - real numbers, written in decimal (0.1234) or scientific notation (1.23e4)
  - numbers are double by default (3 is stored as 3.00)
  - three special doubles: `Inf`, `-Inf`, and `NaN` (not a number)
- **integer** - integers, whole numbers followed by `L` (3L or 1e3L)
- **character** - strings with single or double quotes ('hello world!' or "hello world!")
- **logical** - boolean, written (`TRUE` or `FALSE`) or abbreviated (`T` or `F`)
- **complex** - complex numbers, where `i` is the imaginary number (5 + 3i)
- **raw** - stores raw bytes

To create atomic vectors:

- `c(2,4,6)` - `c()` function for combining elements, returns 2 4 6
- `2:4` - `:` notation to construct a sequence of integers, returns 2 3 4
- `seq(from = 2, to = 6, by=2)` - `seq()` function to create an evenly-spaced sequence, returns 2 4 6

To check an object's data type:

- `typeof(x)` - returns the data type of object x
- `is.*(x)` - test if object x is data type, returns `TRUE` or `FALSE`
  - `is.double()`
  - `is.integer()`
  - `is.character()`
  - `is.logical()`

To change an object to data type (**explicit coercion**):

- `as.*(x)` - coerce object to data type

- `as.double()`
- `as.integer()`
- `as.character()`
- `as.logical()`

Note that atomic vectors must contain only elements of the same type. If you try to include elements of different types, R will coerce them into the same type with no warning (**implicit coercion**) according to the heirarchy `character > double > integer > logical`.

## Operations

Arithmetic operators:

- `+` - add
- `-` - subtract
- `*` - multiply
- `/` - divide
- `^` - exponent

Comparison operators return true or false:

- `a == b` - equal to
- `a != b` - not equal to
- `a > b` - greater than
- `a < b` - less than
- `a >= b` - greater than or equal to
- `a <= b` - less than or equal to

Logical operators combine multiple true or false statements:

- `&` - and
- `|` - or
- `!` - not
- `any()` - returns true if any element meets condition
- `all()` - returns true if all elements meet condition
- `%in%` - returns true if any element is in the following vector

Most math operations (and many functions) are **vectorized** in R:

- they can work on entire vectors, without the need for explicit loops or iteration.
- this a powerful feature that allows you to write cleaner, more efficient code
- To illustrate, suppose `x <- c(1, 2, 3)`:
  - `x + 100` returns `c(101, 102, 103)`
  - `x == 1` returns `c(TRUE, FALSE, FALSE)`



## More complex structures

Some more complex data structures are **built from atomic vectors** by adding **attributes**:

- **matrix** - a vector with a **dim** attribute representing 2 dimensions
- **array** - a vector with a **dim** attribute representing n dimensions
- **factor** - an integer vector with two attributes: **class="factor"** and **levels**, which defines the set of allowed values (useful for categorical data)
- **date-time** - a double vector where the value is the number of seconds since Jan 01, 1970 and a **tzzone** attribute representing the time zone
- **data.frame** - a named list of vectors (of equal length) with attributes for **names** (column names), **row.names**, and **class="data.frame"** (used to represent datasets)

To create more complex structures:

- **list(x=c(1,2,3), y=c('a','b'))** - create a list
- **matrix(x, nrow=2, ncol=2)** - create a matrix from a vector **x** with **nrow** and **ncol**
- **array(x, dim=c(2,3,2))** - create an array from a vector **x** with dimensions
- **factor(x, levels=unique(x))** - turn a vector **x** into a factor
- **data.frame(x=c(1,2,3), y=c('a','b','c'))** - create a data frame

Missing elements and empty vectors:

- **NA** - used to represent missing or unknown elements in vectors. Note that **NA** is contagious: expressions including **NA** usually return **NA**. Check for **NA** values with **is.na()**.
- **NULL** - used to represent an empty or absent vector of arbitrary type. **NULL** is its own special type and always has length zero and **NULL** attributes. Check for **NULL** values with **is.null()**.

## Built-in functions

Note that you do not need to memorize these built-in functions to be successful on quizzes. Use this as a reference.

For basic math:

- **log(x)** - natural log
- **exp(x)** - exponential
- **sqrt(x)** - square root
- **abs(x)** - absolute value
- **max(x)** - largest element
- **min(x)** - smallest element
- **round(x, n)** - round to n decimal places
- **signif(x, n)** - round to n significant figures

- `sum(x)` - add all elements

For stats:

- `mean(x)` - mean
- `median(x)` - median
- `sd(x)` - standard deviation
- `var(x)` - variance
- `quantile(x)` - percentage quantiles
- `rank(x)` - rank of elements
- `cor(x, y)` - correlation
- `lm(x ~ y, data=df)` - fit a linear model
- `glm(x ~ y, data=df)` - fit a generalized linear model
- `summary(x)` - get more detailed information from a fitted model
- `aov(x)` - analysis of variance

For vectors:

- `sort(x)` - return sorted vector
- `table(x)` - see counts of values in a vector
- `rev(x)` - return reversed vector
- `unique(x)` - return unique values in a vector
- `array(x, dim)` - transform vector into n-dimensional array

For matrices:

- `t(m)` - transpose matrix
- `m %*% n` - matrix multiplication
- `solve(m, n)` - find  $x$  in  $m * x = n$

For data frames:

- `view(df)` - see the full data frame
- `head(df)` - see the first 6 rows of data frame
- `nrow(df)` - number of rows in a data frame
- `ncol(df)` - number of columns in a data frame
- `dim(df)` - number of rows and columns in a data frame
- `cbind(df1, df2)` - bind dataframe columns
- `rbind(df1, df2)` - bind dataframe rows

For strings:

- `paste(x, y, sep=' ')` - join vectors together element-wise
- `toupper(x)` - convert to uppercase
- `tolower(x)` - convert to lowercase

- `nchar(x)` - number of characters in a string

For simple plotting:

- `plot(x)` values of x in order
- `plot(x, y)` - values of x against y
- `hist(x)` - histogram of x

## Programming in R

Writing **functions** and handling **control flow** are important aspects of learning to program in any language. For our purposes, some general conceptual knowledge on these topics is sufficient (see below). Those interested to learn more might enjoy the book [Hands-On Programming with R](#).

- **Functions** are reusable pieces of code that take some input, perform some task or computation, and return an output.

```
function(inputs){  
  ## do something  
  return(output)  
}
```

- **Control flow** refers to managing the order in which expressions are executed in a program:
  - **if...else** - if something is true, do this; otherwise do that
  - **for** loops - repeat code a specific number of times
  - **while** loops - repeat code as long as certain conditions are true
  - **break** - exit a loop early
  - **next** - skip to next iteration in a loop

## Further reading and references

Suggested further reading:

- [Getting started with Data in R](#) from Modern Dive textbook
- [R Nuts and Bolts](#) in R Programming for Data Science by Roger Peng
- [Base R Cheat Sheet](#)

Other references:

- [Vectors](#) in Advanced R by Hadley Wickham
- [A field guide to base R](#) in R for Data Science by Hadley Wickham